



Security Copilot Defender

- Attendee Lab Manual

Table of Contents

Purpose of This Lab	4
What You Will Learn	4
Prerequisites	4
Accounts	4
Environment	4
Important Notes.....	4
Activity 0: Preparation Activity: Phishing Triage Agent	5
Activity 1: Microsoft Defender XDR	6
Exercise 1: Embedded Incident Investigation	6
Open the Incident.....	6
Generate the Incident Report	10
File Analyzer.....	10
Script Analyzer	13
Device Summary	16
Exercise 2: Threat Hunting Agent.....	18
Open Advanced Hunting.....	18
Run a Natural-Language Hunt	19
Pivot to a Device	20
Expand to a User.....	21
Summarize Your Hunt.....	22
Exercise 3: Phishing Triage Agent.....	22
Objective.....	23
Open the Incident Queue	23
Review How the Agent Processes Submissions	24
(View Only): How Feedback Works	32
Monitor Agent Performance	34
Exercise 4: Microsoft Security Copilot Dynamic Threat Detection Agent (preview)	38
Objective.....	38
Open the Incident Queue	39

Exercise 5: Microsoft Security Copilot Threat Intelligence Briefing Agent in Microsoft Defender ..	43
Objective.....	43
Exercise 6: Standard operation procedure (SOP) for Guided Response	46
Description	46

Purpose of This Lab

This hands-on lab introduces Microsoft Security Copilot embedded in Microsoft Defender XDR.

You will investigate incidents, run threat hunts, explore agents, and learn how generative AI accelerates SOC workflows.

What You Will Learn

- How to use Security Copilot within Defender XDR
- How to perform end-to-end incident investigation
- How to leverage the Threat Hunting Agent, Phishing Triage Agent, and Conditional Access Agent
- How to save prompts into reusable promptbooks

Prerequisites

To complete the lab successfully, you need:

Accounts

- A valid [Zava Alpine Ski House account](#)
- A separate **browser profile** to avoid authentication conflicts

Environment

- Microsoft Defender portal access
- Basic understanding of incidents, alerts, and identity concepts

Important Notes

Shared Environment

Incidents, alerts, users, and data change frequently.

Your results may differ from screenshots shown in the lab.

Do Not Perform Remediation Actions


Avoid isolating devices, disabling users, deleting apps, or making policy updates.

Activity 0: Preparation Activity: Phishing Triage Agent

The Phishing Triage Agent only triggers when a user manually reports a message as phishing. This cannot be automated.

Before the workshop, each participant should:

1. Go to **<https://outlook.office.com>**
2. Select any email
3. Choose **Report** → **Phishing**




Report phishing

Reporting phishing scams helps keep your and other people's personal information safer.

☐ Don't show me this message again



This action creates a user submitted phishing incident that the Phishing Triage Agent will automatically classify-



Activity 1: Microsoft Defender XDR

This activity teaches you how to investigate incidents using **Microsoft Defender XDR** with embedded **Security Copilot** capabilities.-by-step analysis across incidents, files, scripts, devices, and ultimately use Security Copilot to summarize and streamline your investigation.


You will perform stepbystep analysis across incidents, files, scripts, devices, and ultimately use Security Copilot to summarize and streamline your investigation.

Exercise 1: Embedded Incident Investigation

Open the Incident

1. In the **Microsoft Defender portal**, go to:
Investigation & Response → **Incidents & Alerts** → **Incidents**
2. Set the date filter to **Last 6 months**.
3. Search for and open the **HumanOperated Ransomware** incident. For example: **Incident ID: 36193** [Incident - Microsoft Defender](#) (Open this in a browser where you are logged in using your lab credentials)

💡 Opening an incident automatically triggers Copilot to generate an Incident Summary on the right side.




Review of the Copilot Incident Summary

- Read the **automatically generated summary**.
- Identify:
 - Key entities
 - Attack story
 - Initial impact

💡 *This gives analysts a fast, high-quality understanding without reading the full activity log.*

[Summarize incidents with Microsoft Copilot in Microsoft Defender - Microsoft Defender XDR | Microsoft Learn](#)



Explore Guided Response Cards

Scroll through the **Guided Response** cards in the Copilot pane.

Each card shows:

- Recommended action.

💡 *Copilot provides containment actions like “isolate device” or “disable user” to stop attacks immediately.*

- Entity affected
- Reason for the recommendation
- Whether the action has been previously completed

- Home
- Exposure management
- Investigation & response
- Threat intelligence
- Assets
- Microsoft Sentinel
- Identities
- Endpoints
- Email & collaboration
- Cloud apps
- Cloud security
- Operational technology
- Security Copilot
- Cases
- SOC optimization

Incidents - 7 - Muman operated ransomware attack was launched from a compromised asset (Attack disruption)

ID 36193: Human-operated ransomware attack was launched from a compromised asset (Attack disruption)

High
Active
Mimik Evals - AlpineSkiHouse
Intrusion alert
Last update time: Jan 21, 2026 1:50 AM

Ransomware
Critical asset
Credential Phish
Lateral Movement
Attack Disruption
Predictive Shielding
HumOR
LATEST
+2

Attention: Attack disruption initiated multiple response actions. For more details, go to the [Activities tab](#) or visit the [Action center](#).

Go Hunt queries launched from the entity menu now default to a time range starting from the incident's start time up to the execution time. For incidents older than 30 days, the query defaults to last 30-day.

Select Blast radius from an entity in the incident graph to identify any possible attack paths from that entity. [Learn about the blast radius and graph visualizations](#)

Attack story | Alerts (44) | Activities (205) | Assets (13) | Investigations (0) | Evidence and Response (81) | Recommended actions (24) | Summary

Detection & Categories

Active alerts	Categories
0	5

First activity

Dec 2, 2025 2:48:30 PM

Last activity

Dec 2, 2025 7:46:51 PM

Creation time

Dec 2, 2025 4:57:32 PM

Alerts

Dec 2, 2025 2:48 PM
Resolved
Compromised account conducting hands-on-keyboard attack
rayt-pc.zava-corp.com
rayt

Report incident history

Incident graph

Layout | Group similar nodes

Communication | Association

Priority assessment

100

This incident is ranked as top priority and requires immediate attention.

Notable priority factors:

- 3 Notable alert types
 - Azure VM extension activity followed by ransomware or hands-on-keyboard attack.
 - Possible pass-the-hash authentication.
 - Potentially compromised assets exhibiting ransomware-like behavior.
- 1 High risk threats
 - Credential Phish
- 1 Critical tags involved

Tasks

AI-generated content may be incorrect. Check it for accuracy.

New
Copilot

Isolate device jump-srv.zava-corp.com

Other organizations tend to take this action for similar incidents.

Isolate device
View device

AI-generated content may be incorrect. Check it for accuracy.

New
Copilot

Delete similar emails

We found emails that are very similar to emails involved in this incident. Delete these emails to contain the attack.

Soft delete emails
View similar emails

AI-generated content may be incorrect. Check it for accuracy.

New
Copilot

Suspend the account johndoe

Other organizations tend to take this action for similar incidents.

Disable
View user

AI-generated content may be incorrect. Check it for accuracy.

New
Copilot


Suspend the account 27bf9e71-106d-49bd-

Page | 9



Generate the Incident Report

1. At the top of the Copilot side panel, select the **Report icon** (or use ... → **Generate incident report** on the incident page).



Tip:

This report streamlines what traditionally takes analysts 20–40 minutes of manual triage and combines alerts, logs, scripts, and device data into a single, defensible narrative.

Once generated, the report can be copied to your preferred system.

File Analyzer


1. Access the **Attack Story Graph tab**: Click on the Files option in the Attack Story Graph.
2. View Files: Click on **View <nr> Files** to see the list of files associated with the incident.
3. **Expand Node Details**: To get more information about a specific file, expand the Node Details.




Tip:

Alternatively, choose a file to investigate in the alert page.


Security Copilot – Lab Manual Last Changed: 0204



4. Select the file **Mimikatz.exe** file to investigate.



5. Then select **Open File Page**. Copilot automatically generates an analysis upon opening a file page.
6. Click on **Show details**.



The screenshot displays the Microsoft Defender Security Center interface. The left sidebar contains navigation options such as Home, Exposure management, Investigation & response, Threat intelligence, Assets, Microsoft Sentinel, Identities, Endpoints, Email & collaboration, Cloud apps, Cloud security, Operational technology, Security Copilot, Cases, SOC optimization, and Reports. The main content area shows the file details for 'mimikatz.exe', which is an unsigned file of 1.36 MB. The file details section includes SHA1, SHA256, and MD5 hashes, file size, and PE type. The analysis results section shows 'No active alerts or incidents' and 'No device prevalence'. The VirusTotal section indicates a '66/72 (High risk)' score. The Malware detection section shows '1 malware type found'. The right sidebar displays the Copilot-generated file analysis, which includes an overview of the file's prevalence and a warning about its authenticity.

Note: Please note that screenshots may differ between the lab and demo environments

💡 The file analysis results generated by Copilot typically includes assessment of the file, detection name if malicious/unwanted, key file information, Highlights strings found, lists API calls used, and relevant certificate information.



Script Analyzer

Return to the Incident

1. At the top of the page, click the **incident title** to go back to the full incident view.

Incidents > Human-operated ransomware attack was launched from a compromised asset (attack disruption)


Open the Script Analyzer

2. In the left alert list, select **Possible attempt to access Primary Refresh Token (PRT)**.
Select time stamp **:39 PM (In this case 4:39 PM)


The screenshot displays the Microsoft Sentinel console interface. On the left, a navigation pane lists various security tools and categories. The main area shows an incident titled 'ID 36193: Human-operated ransomware attack was launched from a compromised ...'. Below the incident title, there's a section for 'Alerts' with a list of alerts. One alert, 'Possible attempt to access Primary Refresh Token (PRT)', is highlighted with a red box. To the right of the alert list, there's a section for 'Incident graph' and 'Process tree'. The 'Process tree' section shows a sequence of processes: 'userinit.exe', 'explorer.exe', and 'powershell.exe'. The 'powershell.exe' process is expanded, showing a command that includes a 'curl' command. The 'Alert timeline' section shows the sequence of events, including 'userinit.exe', 'explorer.exe', and 'powershell.exe'.

3. In the alert details, find the **PowerShell script** that includes the **curl command**.
Select Process ID 12868 in PowerShell where the curl command is demonstrated at **:37:52 PM.
- 4.

Security Copilot – Lab Manual Last Changed: 0204



5. Select the script and click **Analyze** to start Script Analyzer.



Tip:

The PRT-access script (curl executed via PowerShell) is an excellent example for Script Analyzer. If you can't find it, analyzing **any script or command** will still show you how Script Analyzer works.

Security Copilot – Lab Manual Last Changed: 0204

Attack story Alerts (71) Assets (13) Investigations (4) Evidence and Response (89) Recommended actions (18) Summary ...

Alerts

Play attack story Unpin all Show all

- Jun 20, 2024 10:51 AM** Resolved
Possible attempt to access Primary Refresh Token (PRT)
vnevado-win10v.vnevado.alpineskihouse.co jonaw
- Jun 20, 2024 10:58 AM** Resolved
Mimikatz credential theft tool
vnevado-win10v.vnevado.alpineskihouse.co jonaw
- Jun 20, 2024 10:58 AM** Resolved
Mimikatz credential theft tool
vnevado-win10v.vnevado.alpineskihouse.co
- Jun 20, 2024 10:59 AM** Resolved
Unknown malware was detected and was active
vnevado-win10v.vnevado.alpineskihouse.co

Incident graph

Layout Group similar nodes

Communication Association

Possible attempt to access Primary Refresh...

[2484] powershell.exe -Enco... Remote execution

Command line Analyze

Command line (decoded) Analyze

curl vectorsandarrows.com

function Get-UserPRTToken



Device Summary

Within the incident page, select the **Assets** tab.

Incidents > Human-operated ransomware attack was launched from a compromised asset (att

ID 36193: Human-operated ransomware attack

■■■ High | ● Active | 👤 Mimik Emails - AlpineSkiHouse | ✎ True alert | ⌚ Last update

Ransomware Critical asset Credential Phish Lateral Movement Attack Disruption Predictive Sh

Sorry, we're having trouble loading the page. Try refreshing the page later. (page: Inci

Attack story Alerts (44) Activities (205) **Assets (13)** Investigations (0)

All assets (13)

- Devices (5)**
- Users (6)
- Mailboxes (0)
- Apps (0)
- Cloud Resources (2)

Devices

↓ Export ⌘ Manage tags ⬡ Initiate au

☐	Device name ▾
☐	rayt-pc-vm
☐	jump-srv.zava-corp.com
☒	rayt-pc.zava-corp.com
☐	laurenceg-pc.zava-corp.com
☐	lydiab-pc.zava-corp.com

Click **Devices** and **select a device on the right**.

Click **Open Device** page. Copilot automatically generates an analysis upon opening the device page.

The screenshot displays the Microsoft Defender Security Center interface. On the left is a navigation pane with various security tools. The main area shows an incident titled 'ID 36193: Human-operated ransomware attack'. Below this, there's a section for 'All assets (13)' with a sub-section 'Devices (5)' highlighted. A list of devices is shown, with 'rayt-pc.zava-corp.com' selected. To the right, the 'rayt-pc' device page is open, showing details like 'Logged on users', 'VM details', and 'Health state'. On the far right, a Copilot window provides a 'Device summary' and a list of 'Vulnerable software' with associated risk levels.

Incident Details: ID 36193: Human-operated ransomware attack. Status: High, Active. Tags: Ransomware, Critical asset, Credential Phish, Lateral Movement, Attack Disruption, Predictive Shield.

Assets (13): Devices (5) highlighted. List includes: rayt-pc-vm, jump-srv.zava-corp.com, rayt-pc.zava-corp.com (selected), laurenceg-pc.zava-corp.com, lydiab-pc.zava-corp.com.

rayt-pc Device Details:

- Category: Computers and Mobile
- Type: Unknown
- Subtype: Workstation
- Discovery sources: Active
- Primary user: Ray Tanaka
- Domain: zava-corp.com
- OS: Windows 10 64-bit (Release 22H2 Build 19045.6691)
- SAM name: RAYT-PC\$
- Health state: Active
- Created on: Jul 31, 2025 2:18:21 PM
- Data sensitivity: None
- IP addresses: Private: 10.2.0.132, Public: 4.152.23.243

Copilot Analysis:

- Device summary:** Feb 4, 2026 6:17 PM. The device "rayt-pc.zava-corp.com" is a Virtual Machine owned by Ray Tanaka (rayt@zava-corp.com). It currently has a high exposure level, indicating significant vulnerability. Multiple critical and high-severity software vulnerabilities are present, and several key security settings are disabled or misconfigured, increasing the risk of compromise.
- Vulnerable software:**
 - openssl 3.5.0.0: Critical — Exposes the device to severe cryptographic vulnerabilities.
 - firefox 127.0.2.0: Critical — Contains critical browser vulnerabilities.
 - windows_10 10.0.19045.6691: High — Operating system is susceptible to high-severity exploits.
 - chrome 143.0.7499.193: High — High-risk.



Tip:
Notice the owner, software vulnerability and risky settings of the device




Exercise 2: Threat Hunting Agent

In this exercise, you will see how Security Copilot assists analysts during threat hunting in Microsoft Defender XDR by generating KQL, identifying anomalies, and guiding pivots using natural language.

⚠️ *This is a shared environment. Query results will vary.*

Open Advanced Hunting


1. Go to **Investigation & response** -> **Hunting** → **Advanced hunting** in the Defender portal.



2. The Security Copilot pane opens automatically. If it does not, please click the Copilot button at the top



3. In the Copilot section, click the three dots (...) and ensure the Threat Hunting Agent is toggled on.



Run a Natural-Language Hunt

4. In the prompt bar, type:

Find suspicious PowerShell activity in the last 90 days.

Copilot will:

- Generate the KQL
- Run the hunt
- Interpret results (e.g., encoded commands, unusual parent processes)




Tip:
Compare the KQL to what you might manually write.

A screenshot of the Microsoft Security Copilot interface. The left sidebar shows the 'Advanced hunting' section selected. The main area displays a query generated by Copilot:

```
1 let suspicious_keywords = dynamic([
46   is_suspicious_cmd,
47   is_suspicious_parent,
48   is_network
49 ])
50 | order by Timestamp desc
```

 Below the query is a table of results with columns: Timestamp, DeviceName, DeviceId, AccountName, and AccountDomain. The table contains several rows of data, including entries for 'vnevado-win11u.vnevad' and 'ssetest'. On the right side, a Copilot chat window is open, showing the prompt 'Find suspicious PowerShell activity in the last 90 days' and a summary of the results. The summary states: '2 step(s) completed. Here's a summary of suspicious PowerShell activity detected in your environment over the last 90 days:'. It lists several findings, including 'Numerous PowerShell executions were flagged as suspicious due to the use of encoded commands, execution policy bypasses, hidden or non-interactive modes, and unusual parent processes' and 'Many events are related to system management or automation tools (e.g., Intune, Microsoft Dependency Agent, HealthScripts)'. At the bottom of the chat window, there is a prompt bar with the text 'Ask a question about your data. A query will be'.

Security Copilot – Lab Manual Last Changed: 0204



Pivot to a Device

From the results, choose a device and ask:

Which device shows recent authentication events ?

Observe how Copilot:

- Automatically adapts the query
- Highlights anomalies such as failed logons or blank accounts

Security Copilot – Lab Manual Last Changed: 0204

Advanced hunting

New query

Schema Functions

Search

Favorites

Alerts & behaviors

Apps & identities

Query

Run query

Set in query

Save

Share link

Create summary rule

Create detection rule

On February 25, 2026, we're changing how advanced hunting parses Boolean values. Learn more

SecurityEvent

where TimeGenerated > ago(7d)

where isNotEmpty(DeviceId)

summarize EventCount = count(), MostRecentAuthEvent = max(TimeGenerated) by DeviceId, DeviceName

sort by MostRecentAuthEvent desc

Results visualization

Group by: DeviceName

Summary by Copilot

Observations

3 main device names show recent authentication events

karinb-pc.zava-corp.com has the highest event count (8)

Most recent activity: VNEVADO-Win11Tvnevado.alpineskhouse.co

Feb 4, 2026 6:35 PM

Which device shows recent authentication events?

Several devices have shown recent authentication events in the last 7 days. The most active devices are:

- VNEVADO-Win11Tvnevado.alpineskhouse.co (DeviceId: SWD/PRINTENUM(D59B6682-B571-4E0D-A6A9-8810813BA298) with 7 authentication events, most recently on 2026-02-04.
- VNEVADO-Win11Tvnevado.alpineskhouse.co and VNEVADO-Win10S.vnevado.alpineskhouse.co also show multiple recent authentication events.
- karinb-pc.zava-corp.com has several device IDs with authentication activity, the most recent being on 2026-01-30.

Would you like to investigate authentication details for a specific device from this list?

See the logic behind the query

Drill into failed authentication attempt

Correlate with recent PowerShell activity

Check for unusual login time

Ask a question about your data. A query will be

Expand to a User

Ask:

Hunt for alerts or unusual activity for involved users in the last 7 days.

Copilot correlates activity across alerts, devices, and identities.

Advanced hunting

New query

Schema Functions

Search

Favorites

Alerts & behaviors

Apps & identities

Query

Run query

Set in query

Save

Share link

Create summary rule

Create detection rule

On February 25, 2026, we're changing how advanced hunting parses Boolean values. Learn more

let device_list = dynamic([AlertSeverity, AlertId, Timestamp, AdditionalInfo])

Results

14 items

Filters

DeviceName

UserName

AlertCount

AlertNames

ActivityTypes

Feb 4, 2026 6:37 PM

Hunt for alerts or unusual activity for involved users in the last 7 days.

3 step(s) completed

Here's what I found for alerts and unusual activity involving users who recently authenticated to the devices you asked about in the last 7 days:

- High volumes of suspicious process executions (PowerShell, cmd.exe, rundll32.exe, schtasks.exe, etc.) were observed on all four devices, mostly under the SYSTEM and LOCAL SERVICE accounts, but also under named users like karinb.
- The activity is flagged as "Medium" severity and is consistent with both legitimate automation (e.g., Intune, Defender, system management) and potentially risky behavior (e.g., encoded PowerShell, fileless script execution, scheduled task creation).
- On karinb-pc.zava-corp.com, user karinb executed PowerShell with obfuscated or fileless payloads, and there were also 4 "Suspicious scheduled task" alerts—a common persistence technique for attackers.
- No high-severity alerts were found, but the volume and nature of these events warrant further review, especially for accounts running encoded or fileless PowerShell.

Would you like to drill down into the details of the suspicious scheduled task alerts on karinb-pc.zava-corp.com for deeper investigation?

Ask a question about your data. A query will be

Summarize Your Hunt

Ask:

Summarize the key findings and recommend next steps.

Expected output includes:

- Top anomalies
- Affected entities
- Recommended investigation actions

The screenshot displays the Microsoft Defender Advanced Hunting interface. On the left is a navigation pane with categories like Home, Exposure management, Investigation & response, Hunting, Actions & submissions, Partner catalog, Threat intelligence, Assets, Microsoft Sentinel, Identities, Endpoints, Email & collaboration, Cloud apps, and Cloud security. The main area is titled 'Advanced hunting' and contains a 'New query' button, a search bar, and a list of favorites. A query is entered in the 'Query' field:

```
57 let device_list = dynamic([
58   "VMEVAD0-WIN105.vnevado.alpineskihouse.co",
59   "karinb-pc.zava-corp.com"
60 ]);
61
62
63 let auth_users = DeviceLogonEvents
64 | where Timestamp > ago(7d)
65 | where DeviceName in~ (device_list)
66 | where IsNotEmpty(AccountName)
67 | project DeviceName, AccountName, AccountDomain, AccountSid, Timestamp
68 | distinct DeviceName, AccountName, AccountDomain, AccountSid;
```

 Below the query, the 'Results' tab shows a 'Summary by Copilot' section. This section includes 'Observations' (Most suspicious activity on 2-3 devices, karinb-pc.zava-corp.com triggered multiple alerts, PowerShell misuse concentrated on specific endpoints) and a 'Results visualization' (a pie chart showing the distribution of results by device name). On the right, a 'Copilot' panel is open, displaying a chat interface with a message: 'Feb 4, 2025 6:39 PM Summarize the key findings and recommend next steps.' The Copilot response includes 'Key Findings' (Suspicious PowerShell Activity, Recent Authentication Events, Alerts & Unusual Activity) and 'Recommended Next Steps' (Investigate Scheduled Task Alerts, Analyze PowerShell Command Line, Examine).

Exercise 3: Phishing Triage Agent

Objective

In this exercise, you will explore how the Phishing Triage Agent automatically analyzes user reported phishing submissions, classifies incidents, provides transparent reasoning, and adapts to your organization through feedback. This gives analysts a faster, more consistent, and more explainable phishing triage workflow.

⚠ Important: Before starting this exercise, ensure you have completed the steps in Preparation Activity: Phishing Triage Agent at the beginning of this lab.

Open the Incident Queue

1. In the **Microsoft Defender portal**, go to:
Investigation & Response → Incidents & Alerts → Incidents.
2. Filter for **user submitted phishing** incidents.
3. In the **Incidents** list, open the **Filters** pane.
4. Set **Last update time** to **Last week** (top-right filter control).
5. Select **Add filter** → **Tags** → choose **Agent**.

💡 This filters the view to only incidents where an agent, such as the Phishing Triage Agent has taken action.

The screenshot shows the Microsoft Defender portal's 'Incidents' page. The left sidebar has 'Incidents & alerts' selected. The main area shows a list of incidents filtered by 'Last update time: 1 Week' and 'Tags: Agent'. The table lists incidents with columns for Incident name, Incident ID, Priority, Tags, Severity, Investigation state, Categories, Impacted assets, and Active alerts.

Incident name	Incident ID	Priority	Tags	Severity	Investigation state	Categories	Impacted assets	Active alerts
Initial access incident involving multiple users	33617	51	Credential Phish	High	2 investigation states	Initial access	2 Accounts, 2 Mailboxes	0/4
Email reported by user as malware or phish involve...	42379	20	Credential Phish	Low	No threats found	Initial access	Kumar Lepp	0/1
Email reported by user as malware or phish involve...	43998	18	Credential Phish	Low	No threats found	Initial access	w17426	0/1
Email reported by user as malware or phish	43941	18	Credential Phish	Low	No threats found	Initial access	Lennart Hlatshwayo	0/1
Email reported by user as malware or phish	43933	18	Credential Phish	Low	Terminated by system	Initial access	Lennart Hlatshwayo	0/1
Email reported by user as malware or phish	43877	18	Credential Phish	Low	Queued	Initial access	w17426	0/1

6. From the filtered list, pick **any incident named:**
Email reported by user as malware or phish

💡 The Phishing Triage Agent operates on user submitted phish alerts and begins classifying submissions as soon as it is enabled.-submitted phish alerts and begins classifying submissions as soon as it is enabled.

Review How the Agent Processes Submissions

1. Select a phishing submission incident. On the incident page, look for the **Phishing Triage Agent** card in the Copilot/Tasks side panel under **Guided Response** → **Triage**. [Incident - Microsoft Defender](#) (Incident ID: 41032)

The screenshot displays the Microsoft Defender Security Center interface for incident ID 41032. The title bar reads "ID 41032: Email reported by user as malware or phish involving one user". The interface is divided into several sections:

- Left Sidebar:** Contains navigation links for Home, Exposure management, Investigation & response, Threat intelligence, Assets, Microsoft Sentinel, Identities, Endpoints, Email & collaboration, Cloud apps, Cloud security, Operational technology, Security Copilot, Cases, SOC optimization, Reports, Learning hub, Trials, More resources, System, and Customize navigation.
- Central Incident Graph:** Shows a graph with nodes for "u6251ba.alpineshouse.com" and "235.253.253.253", connected by a line labeled "Interesting news".
- Right Sidebar:** Contains several task panels:
 - Completed:** A panel titled "Classify phishing email attempt as 'True positive' for alert ID fa3f312baa-44d9-46fa-02da-08a014e09192". It contains a detailed analysis of the phishing attempt, noting that the email is classified as malicious due to a suspicious sender address and a link to an unrelated, potentially dangerous domain. The investigation found that the sender's address was inconsistent with the display name, which is a common phishing indicator. The email message was generic and included a link leading to a domain unrelated to the supposed sender, increasing the likelihood of malicious intent. There was no evidence of prior communication or any indication that this was a phishing simulation. The overall structure and content of the email matched known phishing tactics, with no strong evidence suggesting a benign purpose.
 - Threat Overview Profile:** A panel titled "Evolving phishing threats" showing 81 impacted assets and a "View threat analytics report" button.
 - Contain:** A panel with a "New" button and a "Custom guidebook" link. It includes a "Submit email to Microsoft" button and a "Mark as completed" button.
 - Block malicious URLs:** A panel with a "New" button and a "Custom guidebook" link. It includes a "Block all malicious URLs identified in the phishing campaign" button.

 **Completed**

 **Phishing Triage Agent** 

Classify phishing email attempt as 'True positive' for alert ID fa3f312baa-44d9-a6fa-02da-08de51e09592

This email is classified as malicious because it contains a suspicious sender address that does not match the display name and includes a link to an unrelated, potentially dangerous domain. The investigation found that the sender's address was inconsistent with the display name, which is a common phishing indicator. The email message was generic and included a link leading to a domain unrelated to the supposed sender, increasing the likelihood of malicious intent. There was no evidence of prior communication or any indication that this was a phishing simulation. The overall structure and content of the email matched known phishing tactics, with no strong evidence suggesting a benign purpose.

[See less](#)

Change classification

[View agent activity](#)

AI-generated content may be incorrect. Check it for accuracy.

2. Observe:
 - **Verdict assigned** (True Positive or False Positive)
 - **Actions taken**, such as resolving a false positive
 - **Reasoning summary** explaining *why* the agent reached this decision
3. You can select the **More actions** ellipsis to view more alert details, copy the agent's classification details to the clipboard, or manage feedback.

✓ Completed

Phishing Triage Agent ...

Classify phishing email at
for alert ID fa3f312baa-44
08de51e09592

Copy to clipboard

View alert details

Manage feedback

This email is classified as n
contains a suspicious sender address that does not
match the display name and includes a link to an
unrelated, potentially dangerous domain. The
investigation found that the sender's address was
inconsistent with the display name, which is a
common phishing indicator. The email message was
generic and included a link leading to a domain
unrelated to the supposed sender, increasing the
likelihood of malicious intent. There was no
evidence of prior communication or any indication
that this was a phishing simulation. The overall
structure and content of the email matched known
phishing tactics, with no strong evidence
suggesting a benign purpose.

[See less](#)

Change classification

View agent activity

AI-generated content may be incorrect. Check it for accuracy.

 ✓



Email reported by user as malware or phish

Low
Unknown
Resolved

[Open alert page](#)
[Manage alert](#)
[View submission](#)

Alert state

Classification	Assigned to
True positive	Unassigned
Set Classification	

Alert details

Alert ID	Category
fa3f312baa-44d9-a6fa-02da-08de51e09592	Threat management
MITRE ATT&CK Techniques	Detection source
T1566: Phishing	MDO
Service source	Detection technology
Office 365	-
Generated on	First activity
Jan 13, 2026 12:44:35 AM	Jan 13, 2026 12:43:00 AM
Last activity	Workspace
Jan 13, 2026 12:44:00 AM	-

Evidence

Entity Name	Remediation Status	Verdict
(ip) 255.255.255.255		Suspicious

Alert description

This alert is triggered when any email message is reported as malware or phish by users -V1.0.0.3

Alert Policy

Tasks

AI-generated content may be incorrect. Check it for accuracy.

Triage

New Copilot

Contact user u6251@a.alpineskihouse.co on Teams, and ask them to confirm their activity

Hi.

We noticed that on January 12, 2026 at 13:43 UTC, you reported an email as suspicious or potentially

[See more](#)

[Contact user in Teams](#)
[View user](#)

AI-generated content may be incorrect. Check it for accuracy.

Completed

Phishing Triage Agent

Classify phishing email at for alert ID fa3f312baa-44d9-a6fa-02da-08de51e09592

Copy to clipboard

View alert details

Manage feedback

This email is classified as n... contains a suspicious sender address that does not match the display name and includes a link to an unrelated, potentially dangerous domain. The investigation found that the sender's address was inconsistent with the display name, which is a common phishing indicator. The email message was generic and included a link leading to a domain unrelated to the supposed sender, increasing the likelihood of malicious intent. There was no evidence of prior communication or any indication that this was a phishing simulation. The overall structure and content of the email matched known phishing tactics, with no strong evidence suggesting a benign purpose.

[See less](#)

[Change classification](#)
[View agent activity](#)

AI-generated content may be incorrect. Check it for accuracy.

Threat Overview Profile: Evolving phishing threats

81 impacted assets

[View threat analytics report](#)

- To view the steps the agent took prior to reaching its classification, select **View agent activity** in the Phishing Triage Agent card. This shows the logic behind the agent's final classification.

Security Copilot – Lab Manual Last Changed: 0204

The screenshot displays the Microsoft Defender Security Center interface. On the left, a navigation pane lists various security features. The main area is divided into three sections: Overview, Performance, and Activity. The Overview section provides information about the Phishing Triage Agent, including its purpose and configuration. The Performance section shows a table of activity logs. The Activity section displays a list of events, with one event highlighted. On the right, a panel titled 'Email reported by user as malware or phishing' shows details about a specific email, including its classification and the user's report.

Agents > Phishing Triage Agent

Phishing Triage Agent
Active Microsoft

Overview

About this agent
The Microsoft Defender Phishing Triage Agent triages user-submitted phishing incidents. Leveraging AI to autonomously classify and resolve alerts as either truly malicious or as false alarms, the agent provides a transparent rationale for its classification verdicts in natural language, and uses feedback provided by analysts to continuously improve its accuracy. Since user-submitted phishing incidents are often high volume and require significant time to triage, the Phishing Triage Agent frees up time for high-value work.

Trigger
This agent will run automatically when a user reports a suspicious email.

Permissions
Permissions must be managed by unified role-based access control (URBAC) and include:
• Security data basics (read)
• Email & collaboration content (read)
• Email & collaboration metadata (read)
• Security Copilot trends
• Alerts (manage)

Identity
Connected to agent ID: 8f7d8b61-25d5-4d99-b9c2-3d8b9ac42c7

Plugins
• Microsoft Defender XDR
• Microsoft Threat Intelligence
• Phishing Triage Agent

Role-based access
• Security admins can manage this agent after deployment
• All users with at least the same permissions as the agent can view agent output.

Activity
Review the agent's previous runs.

Start time	Method
1/14/2026, 5:04 AM	Automated
1/13/2026, 12:46 AM	Automated
1/13/2026, 12:45 AM	Automated
1/9/2026, 9:48 AM	Automated
1/9/2026, 9:15 AM	Automated
1/8/2026, 11:36 PM	Automated
1/8/2026, 7:49 PM	Automated
12/24/2025, 1:51 AM	Automated
12/19/2025, 10:54 PM	Automated
12/19/2025, 1:50 AM	Automated
12/19/2025, 1:48 AM	Automated
12/17/2025, 1:19 PM	Automated
12/17/2025, 12:52 PM	Automated
12/17/2025, 12:50 PM	Automated
12/17/2025, 2:43 AM	Automated

Email reported by user as malware or phishing
See full agent workflow

Results
This email is classified as malicious because it contains a suspicious sender address that does not match the display name and includes a link to an unrelated, potentially dangerous domain. The investigation found that the sender's address was inconsistent with the display name, which is a common phishing indicator. The email message was generic and included a link leading to a domain unrelated to the supposed sender, increasing the likelihood of malicious intent. There was no evidence of prior communication or any indication that this was a phishing simulation. The overall structure and content of the email matched known phishing tactics, with no strong evidence suggesting a benign purpose.
See less
If generated content may be inaccurate, check it for accuracy.

Details

Run by	Duration
Phishing Triage Agent	42.5 sec

Started
Jan 13, 2026 12:45 AM

Ended
Jan 13, 2026 12:46 AM

Tasks included
13 tasks

Verdict
True positive

Evidence

Entity Name
• Email reported by user as malware or phishing
• Interesting news
• Lennart Hatzway
• 255.255.255.255
• https://d1748785.blacksky-42e0306.eastus.azurecontainerapps.io/

Email reported by user as malware or phish

See full agent workflow

Results

This email is classified as malicious because it contains a suspicious sender address that does not match the display name and includes a link to an unrelated, potentially dangerous domain. The investigation found that the sender's address was inconsistent with the display name, which is a common phishing indicator. The email message was generic and included a link leading to a domain unrelated to the supposed sender, increasing the likelihood of malicious intent. There was no evidence of prior communication or any indication that this was a phishing simulation. The overall structure and content of the email matched known phishing tactics, with no strong evidence suggesting a benign purpose.

See less

AI-generated content may be incorrect. Check it for accuracy.

Details

Run by	Duration
Phishing Triage Agent	42.5 sec
Started	Ended
Jan 13, 2026 12:45 AM	Jan 13, 2026 12:46 AM
Tasks included	Verdict
13 tasks	True positive

Evidence

Entity Name

Email reported by user as malware or phish

Interesting news

Lennart Hlatshwayo



255.255.255.255

https://ld1748785.blacksky-82e03036.eastus.azurecontainerapps.io/

Select full agent workflow



Page | 29

Security Copilot – Lab Manual Last Changed: 0204



Expand Tasks

Security Copilot – Lab Manual Last Changed: 0204




💡 *The agent provides transparent explanations for each verdict, supporting analyst confidence and auditability.*

(View Only): How Feedback Works

Even though you *cannot* submit feedback in this shared lab environment, you should see **where** analysts normally provide it and **how** it is used.

1. In the incident page, scroll to the **agent activity** area.
2. Locate the **Change Classification** option. This is where analysts correct the agent's classification.



3. Feedback is entered in **natural language**, which makes the agent highly adaptive, it learns organizational nuances, user behavior patterns, and phishing norms from plain English explanations.

Manage alert

Status
Resolved

Assign to
Unassigned

Classification
False positive - Not malicious

Why did you change this classification? ⓘ
This sender is our corporate compliance training vendor, so this email is legitimate.

☒ Use this feedback to teach the agent
Note: feedback must first be evaluated.

[Evaluate feedback](#)

Comment
Add comment

By pressing save, you agree that your feedback will be used to improve Microsoft products and services. You'll be sharing the original request or prompt, the output from Copilot, and additional logs.
[Learn about what's collected](#)
[Privacy Statement](#)

[Save](#) [Cancel](#)

-English explanations.


💡 Choose **Evaluate feedback** to allow you to preview how the agent translates your feedback into a lesson and assess whether the outcome aligns with your intent.

4. In this lab: **Do not Save feedback**, simply observe where the option appears and understand how it influences future triage decisions.


💡 The agent reads the explanation, updates its internal patterns, and improves its future decisions through contextual learning.

Monitor Agent Performance

Open the **Phishing Triage Agent card** in the incident queue, select--> Investigation & response → Incidents & alerts → Incidents



Select *the Manage agent* on the card to open the **Phishing Triage Agent** page.



Tip:

You can also access it through the Defender portal left navigation:
Security Copilot → Agents → Phishing Triage Agent.

This page consists of two tabs: **Overview** and **Performance**.

Overview Tab

The **Overview** tab provides details about the agent's status, identity, role, and recent activity.

Security Copilot – Lab Manual Last Changed: 0204

The screenshot shows the 'Phishing Triage Agent' overview page. The left sidebar contains navigation links: Home, Exposure management, Investigation & response, Threat intelligence, Assets, Microsoft Sentinel, Identities, Endpoints, Email & collaboration, Cloud apps, Cloud security, Operational technology, Security Copilot (selected), Cases, SOC optimization, and Reports. The main content area is divided into sections: 'About this agent' (describing the agent's function), 'Trigger' (stating it runs on suspicious email reports), 'Permissions' (listing access to security data, email content, and alerts), and 'Activity' (a table of recent runs). The 'Activity' table has columns for Start time, Method, and Status, showing several completed runs. Red boxes highlight the 'Phishing Triage Agent' header, the 'Overview' tab, the 'About this agent' section, the 'Permissions' section, and the 'Activity' table.

Start time	Method	Status
2/4/2026, 4:42 PM	Automated	Completed
2/4/2026, 3:54 AM	Automated	Completed
2/4/2026, 3:01 AM	Automated	Completed
2/3/2026, 5:22 PM	Automated	Completed
1/23/2026, 1:27 PM	Automated	Completed
1/22/2026, 8:26 PM	Automated	Completed
1/19/2026, 1:32 AM	Automated	Completed
1/19/2026, 12:52 AM	Automated	Completed
1/18/2026, 9:50 PM	Automated	Completed
1/18/2026, 9:38 PM	Automated	Completed
1/18/2026, 11:32 AM	Automated	Completed

Select an activity from the **Recent activity** list to view details about the agent's investigation and the agent in full workflow.

The screenshot shows the 'Email reported by user as malware or phish' activity details page. The left sidebar is the same as the previous screenshot. The main content area is divided into sections: 'About this agent', 'Trigger', 'Permissions', 'Identity', 'Plugins', 'Role-based access', and 'Activity'. The 'Activity' section shows a list of runs, with the selected activity highlighted. The right sidebar contains a 'Results' section with a detailed description of the email, a 'Details' section with metadata (Run by, Started, Ended, Tasks included, Verdict), and an 'Evidence' section with a list of items. Red boxes highlight the 'Phishing Triage Agent' header, the 'Overview' tab, the 'About this agent' section, the 'Permissions' section, the 'Activity' table, the 'Results' section, and the 'Details' section.

Run by	Duration
Phishing Triage Agent	42.5 sec

Started	Ended
Jan 11, 2026, 12:45 AM	Jan 11, 2026, 12:46 AM


Tasks included	Verdict
13 tasks	True positive

Security Copilot – Lab Manual Last Changed: 0204

Agents > Phishing Triage Agent > Email reported by user as malware or phish

Email reported by user as malware or phish

Completed



Results

The investigation determined the email is NonMalicious because all content, sender verification, and technical

[See more](#)

AI-generated content may be incorrect. Check it for accuracy.

Details

Run by

Phishing Triage Agent

Duration

4 min

Started

Jan 13, 2026 7:04 PM

Ended

Jan 13, 2026 7:08 PM

Tasks included

13 tasks


Verdict

False positive

Performance Tab

💡 This provides an at-a-glance view of the agent's ROI and operational health.

The **Performance** tab displays key metrics about the agent's activity over time, including daily activity, mean time to triage (MTTT), and SCU consumption.



Exercise 4: Microsoft Security Copilot Dynamic Threat Detection Agent (preview)

Objective

In this exercise, you will explore how Microsoft Security Copilot includes the Dynamic Threat Detection Agent, an always-on, adaptive backend service that uncovers hidden threats across Defender and Microsoft Sentinel environments.

The Dynamic Threat Detection Agent uses AI to identify gaps and uncover false negatives by correlating alerts, events, anomalies, and threat intelligence. When the agent identifies a gap, it generates a dynamic alert with the full context in the alert details, including natural language explanations, mapped MITRE ATT&CK techniques, and tailored remediation steps.


Important: *The Dynamic Threat Detection Agent runs automatically in the background. When it generates an alert, the alert shows up in your incidents and alerts queues with Security Copilot as the Detection source.*

Open the Incident Queue

In the **Microsoft Defender portal**, go to:


Investigation & Response → Incidents & Alerts → Incidents.

Filter the Service/Detection Sources to Microsoft Defender XDR, select Security Copilot, and click Apply.




Review Incident ID: 41256 [Incident - Microsoft Defender](#)

Select the Alerts tab, then choose the alert with the Detection Source set to Security Copilot. In this case, the alert name is “Unusual eDiscovery search activity (Preview)”.



Alert URL: [Alert - Microsoft Defender](#)



Security Copilot – Lab Manual Last Changed: 0204

 **Alert description by Dynamic Threat Detection Agent**

AI-generated content may be incorrect. Check it for accuracy.

A sequence of uncommon actions was performed in the compliance center, including viewing case lists and updating and starting a search within an eDiscovery case. These activities originated from an atypical location and service provider, suggesting possible collection or reconnaissance of sensitive information using eDiscovery tools. The behavior is notable due to its rarity and the context of the actions.

This alert is AI-generated and must be verified for accuracy.

Recommended actions

Review the full timeline of compliance center activities for this user, focusing on the scope and content of the eDiscovery searches and any data accessed or exported. Correlate with recent authentication and access patterns to determine if the account or session may have been compromised. Investigate whether the search targeted sensitive or high-value mailboxes and assess if any data exfiltration occurred.

This alert is AI-generated and must be verified for accuracy.

[Read less](#)



Review-->Alert Description-->Alert recommendation action



Unusual eDiscovery search activity (Preview)

High | Unknown | New

 This alert was detected by the Dynamic Threat Detection agent. [Learn more](#)

 Manage alert

 Move alert to another incident

 Tune alert



Entity Name ▾

Remediation Status ▾

Verdict ▾

 88.1.144.248

 Suspicious

Alert description

A sequence of uncommon actions was performed in the compliance center, including viewing case lists and updating and starting a search within an eDiscovery case. These activities originated from an atypical location and service provider, suggesting possible collection or reconnaissance of sensitive information using eDiscovery tools. The behavior is notable due to its rarity and the context of the actions.

This alert is AI-generated and must be verified for accuracy.

Alert recommended actions

Review the full timeline of compliance center activities for this user, focusing on the scope and content of the eDiscovery searches and any data accessed or exported. Correlate with recent authentication and access patterns to determine if the account or session may have been compromised. Investigate whether the search targeted sensitive or high-value mailboxes and assess if any data exfiltration occurred.

This alert is AI-generated and must be verified for accuracy.

Incident details

Incident
[Multi-stage incident involving Initial access & Collection involving multiple users reported by multiple sources](#)

Incident severity
High

Active alerts
3/3

Devices
0

Users
2

Mailboxes
1

Apps
0

Correlation reasons

Exercise 5: Microsoft Security Copilot Threat Intelligence Briefing Agent in Microsoft Defender

Objective

In this exercise, you will explore the Microsoft Security Copilot Threat Intelligence Briefing Agent in Microsoft Defender. The agent generates threat intelligence briefings within minutes by analyzing the latest threat actor activity along with internal and external vulnerability information. It helps security teams save time by producing customized, relevant reports that provide CISOs, security managers, and analysts with critical situational awareness and a strong foundation for defensive actions.

The Threat Intelligence Briefing Agent appears as a banner at the top of the Threat analytics page in the Microsoft Defender portal.

In the Microsoft Defender portal, navigate to Threat Intelligence → Threat Analytics

Threat analytics

Executive summary by Threat Intelligence Briefing Agent

Active 2/4/26, 1:18 PM

Executive Summary

During the reporting period, the insurance industry in Switzerland faced a heightened threat landscape, with financially motivated cybercriminals leveraging the IPv6OS infrastructure...

[View full brief](#) [Run agent](#)

As generated content may be incorrect, check it for accuracy.

Activity 380 Actor 274 Core threat 13 Technique 43 Tool 129 Vulnerability 155 OSINT 5

Latest threats

Threat	Category	Alias	Origin country/region	Targeted industries	Alerts	Impacted assets	Threat exposure level	Misconfigured devices	Vulnerable devices	Published	Last updated
Actor Profile: Storm-2502	Actor	Not available	Nigeria	Not available	0 active / 0	Not available	Not available	Not available	Not available	Feb 4, 2026 3:02 AM	Feb 4, 2026 3:02 AM
Activity Profile: Seashell Blizzard...	Activity	Not applicable	Not applicable	Transportation Systems	0 active / 0	Not available	Not available	Not available	Not available	Feb 3, 2026 11:54 AM	Feb 3, 2026 11:54 AM
Activity Profile: Storm-2372 co...	Activity	Not applicable	Not applicable	+	0 active / 0	Not available	Not available	Not available	Not available	Feb 19, 2025 8:14 AM	Feb 3, 2026 8:14 AM
Activity Profile: SVG smuggling...	Activity	Not applicable	Not applicable	Not applicable	0 active / 0	Not available	Not available	Not available	Not available	Jan 31, 2026 1:24 PM	Jan 31, 2026 1:24 PM
Tool Profile: Bloodhound and ...	Tool	Not available	Not applicable	Not applicable	0 active / 0	Not available	Not available	Not available	Not available	Jan 31, 2026 11:24 AM	Jan 31, 2026 11:24 AM

High-impact threats

Threat	Count	Exposure Level
Activity Profile: Cloud apps used in BIC and phishing	137 / 486	High 70-100
Threat Operator Profile: Cloud identity abuse	55 / 812	Medium 30-69
Threat Operator Profile: Looking phishing threats	45 / 316	Low 0-29

Highest exposure threats

Threat	Count	Exposure Level
Activity Profile: Sophisticated actor attacks Europe	26 / 25	High 70-100

Click **View full brief**. This opens a side panel displaying the Threat Briefing Agent report, generated based on the settings configured under Manage agent settings.

The screenshot shows a dark-themed side panel titled "TI Briefing: Feb 4, 2026 1:18 PM". Below the title, it indicates the status as "Completed" and the start time as "Started on Feb 4, 2026 1:18 PM". A "Manage agent" link is visible. The main section is labeled "Agent Results" with a warning: "AI-generated content may be incorrect. Check it for accuracy." Below this, the "Executive Summary" is displayed, followed by a detailed paragraph about the insurance industry in Switzerland. A red box highlights the "Threats targeting the insurance industry within the Switzerland region" section, which includes a date range of "November 06, 2025 - February 04, 2026". Below this, there are two more sections: "RedVDS-Enabled Attacks Impacting Switzerland" and "Insurance Industry Threat Landscape and Espionage".

TI Briefing: Feb 4, 2026 1:18 PM

Completed | Started on Feb 4, 2026 1:18 PM

[Manage agent](#)

Agent Results AI-generated content may be incorrect. Check it for accuracy.

Executive Summary

During the reporting period, the insurance industry in Switzerland faced a heightened threat landscape, with financially motivated cybercriminals leveraging the RedVDS infrastructure to conduct widespread phishing, business email compromise, and credential theft campaigns. Sophisticated groups also targeted the sector for corporate espionage and selective ransomware deployment, employing advanced intrusion and evasion techniques. These trends underscore the sector's exposure to both monetization-driven and espionage-focused attacks, often enabled by scalable criminal infrastructure. To address these risks, organizations should prioritize multi-layered security controls, enhance user awareness, and strengthen credential management and access policies. Proactive monitoring and automation of detection and response processes are recommended to limit the impact of evolving threats.

Threats targeting the insurance industry within the Switzerland region

Date Range: November 06, 2025 - February 04, 2026


RedVDS-Enabled Attacks Impacting Switzerland

As of January 2026, Switzerland has been specifically impacted by operations leveraging the RedVDS infrastructure, a criminal virtual dedicated server (VDS) platform disrupted by Microsoft and law enforcement in early 2026. RedVDS was widely abused by financially motivated threat actors to facilitate business email compromise (BEC), mass phishing, account takeover, and financial fraud. While RedVDS itself does not select victims, its infrastructure was used by multiple criminal groups to target organizations in Switzerland, particularly those in sectors with substantial banking infrastructure. Attackers used RedVDS to host phishing campaigns, send malicious attachments, and conduct credential theft, often impersonating trusted organizations or suppliers. The platform enabled rapid scaling of attacks, use of anonymization techniques, and supported a broad range of digital currencies for payment, making attribution and disruption challenging. The majority of RedVDS-enabled attacks in Switzerland involved social engineering, phishing, and BEC, with a focus on exfiltrating financial data and credentials. [Storm-2470](#), [RedVDS](#)


Insurance Industry Threat Landscape and Espionage

The insurance industry remains a persistent target for sophisticated cyber threat actors, including those conducting ransomware, corporate espionage, and phishing attacks. Notably, the GOLD BLADE group (also known as RedCurl) has been

Click on "Manage Agent"



View configuration for the Threat Intelligence Briefing Agent



Exercise 6: Standard operation procedure (SOP) for Guided Response


Description

Security Copilot now supports custom guidebooks, allowing your organization to embed its unique incident response playbooks directly into the Microsoft XDR investigation experience.


SOC engineers and team leaders can upload internal playbooks and Standard Operating Procedures (SOPs), which Copilot parses into actionable investigation steps. These steps surface as Guided Response tasks within the incident page and are dynamically tailored to the specific incident type.

By aligning AI-driven response workflows with your organization's playbooks, compliance requirements, and security protocols, Security Copilot delivers more context-aware, trusted, and efficient investigations, helping teams respond faster and with greater confidence.

In the Microsoft Defender portal, navigate to System > Settings, and select Copilot in Defender.



Under **General**, select **Custom Guidebooks**, click **Add new guidebook**, and follow the on-screen steps to upload the SOP document.



In the Microsoft Defender portal, go to **Investigation & response** > **Incidents & alerts** > **Incidents**, and search for incident **41032**. [Incident - Microsoft Defender](#)

In the Copilot side pane, review the Custom guidebook action items generated from the previously uploaded SOP (Custom guidebook).

Security Copilot – Lab Manual Last Changed: 0204

The screenshot displays the Microsoft Defender Security Center interface. The main header shows the incident title: **ID 41032: Email reported by user as malware or phish involving one user**. Below the title, there are tabs for **Attack story**, **Alerts (1)**, **Activities (7)**, **Assets (2)**, **Investigations (1)**, **Evidence and Response (1)**, **Recommended actions (20)**, **Summary**, and **Similar incidents (5)**.

The **Attack story** tab is active, showing a timeline of events. The **Incident graph** is visible in the center, showing a network of nodes and connections. The **Incident details** panel on the right provides information about the incident, including its classification, categories, and creation time.

The **Tasks** panel on the right lists several actions that can be taken to respond to the incident, such as **Submit email to Microsoft**, **Block malicious URLs**, **Block sender domain**, and **Remove phishing emails**. Each task has a **Custom guidebook** link and a **Mark as completed** button.

The **Investigate** panel at the bottom right shows the **Investigate user sign-ins** task, which involves querying Microsoft Entra ID for sign-in logs and filtering by user to look for suspicious activity.