


Microsoft Defender for Cloud

Steps	View
Navigate to https://portal.azure.com portal and open Microsoft Defender for Cloud From there, navigate to the Data and security blade Review the page and note it offers a starting point for your Cloud Data and AI assets	

Steps

Scroll down to the **AI closer look** section, click on **services**. This will open the *Cloud Security Explorer* blade.

Note that the query options are pre-populated, hit **Search**.

The outcome is all the AI Services and Platforms in your Cloud environment, you can drill down further if needed.

This solution helps you to have visibility of current and new AI deployments happening in your environment

View

Microsoft Azure

Search resources, services, and docs (G+)

Copilot

u1209@int.zava-corp.co...
ZAVA (ZAVA-CORP.COM)

Home > Microsoft Defender for Cloud | Data and AI security >

Cloud Security Explorer

Share query linkDownload CSV reportGuides & Feedback

What would you like to search?

Query builderTemplates

15 resource types selected

Clear all

Scope: AllSearch

Search

Showing 1 to 24 of 24 records.

Resource type	Resource name	
AWS Bedrock agents	agent-quick-start-88132	View details >
Azure AI Foundry	ai-contosoenterpriseagenthub399975579637	View details >
Search service	buildsearchservice	View details >
Azure AI project	contosoenterpriseagent	View details >
Azure AI Foundry	ContosoEnterpriseAgent1	View details >
Azure AI Foundry	ContosoEnterpriseAgent2	View details >
Azure AI Foundry	ContosoEnterpriseAgent3	View details >
Azure AI hub	contosoenterpriseagenthub	View details >

Steps

Back to the **Data and security blade** and to the **AI closer look** section in the middle pane you'll see **AI threat detection**, click on **View all AI alerts**.

This will open all the alerts already filtered for your AI workloads.

Select any and view details

For our example we are using *"A Jailbreak attempt on your Azure AI model deployment was blocked by Prompt Shields"*

View

Microsoft Azure

Search resources, services, and docs (G+/)

Copilot

u1209@int.zava-corp.co...
ZAVA (ZAVA-CORP.COM)

Home > Microsoft Defender for Cloud | Data and AI security >

Security alerts

Refresh

Change status

Open query

Suppression rules

Security alerts map

Sample alerts

Alerts workbook

Download CSV report

Guides & Feedback

129
Open alerts

126
Active alerts

3
In progress alerts

2
Affected resources

Open alerts by severity
High (22) Medium (105) Low (2)

Search by ID, IP, name, or resource

Subscription == All

Status == Active, In Progress

Severity == Low, Medium, High

Product Component Name == AI, AIServices

Add filter

Severity	Alert name	Affected resource	Resource Group	Activity start time (UTC+11)	Last updated time (UTC+11)
High	YOUR Azure AI resources have been accessed by an a...	mdc-d4ai-aoai-1	MLL-L4AI	20/10/25, 01:39 pm	20/10/25, 01:39 pm
High	Suspected prompt injection using ASCII smuggling d...	ai-contosoenterpriseagenthub399975579637	MDC-Build-S4AI	28/07/25, 06:48 pm	28/07/25, 06:48 pm
High	A malicious URL was shared in one of your AI applic...	mdc-d4ai-aoai-1	MDC-D4AI	24/09/25, 05:46 pm	24/09/25, 05:46 pm
Medium	A Jailbreak attempt on your Azure AI model deploy...	mdc-d4ai-aoai-1	MDC-D4AI	22/09/25, 12:03 pm	22/09/25, 12:09 pm
Medium	A request by a suspicious user agent was sent to you...	ai-contosoenterpriseagenthub399975579637	MDC-Build-S4AI	28/07/25, 01:15 pm	28/07/25, 01:15 pm
Medium	A request by a suspicious user agent was sent to you...	ai-contosoenterpriseagenthub399975579637	MDC-Build-S4AI	28/07/25, 01:15 pm	28/07/25, 01:15 pm
Medium	A request by a suspicious user agent was sent to you...	ai-contosoenterpriseagenthub399975579637	MDC-Build-S4AI	28/07/25, 01:15 pm	28/07/25, 01:15 pm
Medium	A request by a suspicious user agent was sent to you...	ai-contosoenterpriseagenthub399975579637	MDC-Build-S4AI	28/07/25, 01:15 pm	28/07/25, 01:15 pm
Medium	A request by a suspicious user agent was sent to you...	ai-contosoenterpriseagenthub399975579637	MDC-Build-S4AI	28/07/25, 01:15 pm	28/07/25, 01:15 pm
Medium	A request by a suspicious user agent was sent to you...	ai-contosoenterpriseagenthub399975579637	MDC-Build-S4AI	28/07/25, 11:39 am	28/07/25, 11:39 am
Medium	A request by a suspicious user agent was sent to you...	ai-contosoenterpriseagenthub399975579637	MDC-Build-S4AI	28/07/25, 01:15 pm	28/07/25, 01:15 pm
Medium	A request by a suspicious user agent was sent to you...	ai-contosoenterpriseagenthub399975579637	MDC-Build-S4AI	28/07/25, 01:15 pm	28/07/25, 01:15 pm
Medium	A request by a suspicious user agent was sent to you...	ai-contosoenterpriseagenthub399975579637	MDC-Build-S4AI	28/07/25, 01:15 pm	28/07/25, 01:15 pm
Medium	A Jailbreak attempt on your Azure AI model deploy...	ai-contosoenterpriseagenthub399975579637	MDC-Build-S4AI	28/07/25, 05:25 pm	28/07/25, 05:39 pm

< Previous

Page 1 of 4

Next >

A Jailbreak attempt on your Azure AI model deployment was blocked by Prompt Shields

Medium SeverityActive Status22/09/25, 12:03 pm Activity time

Alert description

There was 1 blocked attempt of a Jailbreak attack on model deployment gpt-4o-model-1 on your Azure AI resource mdc-d4ai-aoai-1.

A Jailbreak attack is also known as User Prompt Injection Attack (UPIA). It occurs when a malicious user manipulates the system prompt, and its purpose is to bypass a generative AI's large language model's safeguards in order to exploit sensitive data stores or to interact with privileged functions. Learn more at https://aka.ms/RAI/jailbreak.

The attempts on your model deployment were using direct prompt injection techniques and were blocked by Azure AI Content Safety Prompt Shields. The prompts were not completed. However, to block further malicious attempts by the suspected user and to handle

View full detailsTake action

Steps

With the alert open, review all details provided like:

- 1. Status
- 2. Alert description
- 3. Affected resources
- 4. MITRE ATT&CK mapping
- 5. Alert details
- 6. Take action

View

Microsoft Azure

Search resources, services, and docs (G+)

Copilot

u1209@int.zava-corp.co... ZAVA (ZAVA-CORP.COM)

Home > Microsoft Defender for Cloud | Data and AI security > Security alerts >

Security alert

00cfe8ee-0724-2ad6-3614-5e9e33a3deec

A Jailbreak attempt on your Azure AI model deployment was blocked by Prompt Shields

Medium

Active

22/09/25, 12:03 pm

Severity

Status

Activity time

Alert description

Copy alert JSON

There was 1 blocked attempt of a Jailbreak attack on model deployment gpt-4o-model-1 on your Azure AI resource mdc-d4ai-acai-1.

A Jailbreak attack is also known as User Prompt Injection Attack (UPIA). It occurs when a malicious user manipulates the system prompt, and its purpose is to bypass a generative AI's large language model's safeguards in order to exploit sensitive data stores or to interact with privileged functions. Learn more at <https://aka.ms/RAI/jailbreak>.

The attempts on your model deployment were using direct prompt injection techniques and were blocked by Azure AI Content Safety Prompt Shields. The prompts were not completed. However, to block further malicious attempts by the suspected user and to handle possible undetected prompt injections, we recommend taking immediate action:

- Investigate the user who created the attempts by looking at the source application's history and consider removing their access.
- Consider there may have been undetected successful prompt injections – investigate to validate no sensitive data was revealed by the model, and that no data poisoning took place.

To get detailed information on the prompt injection attempts, refer to the 'Supporting evidence events' section in the Azure Portal.

Last updated time

Alert details

Take action

General information

Model deployment name

gpt-4o-model-1

Model name

gpt-4o

Model version

2024-08-06

Operation name

ChatCompletions_Create

Investigation steps

Investigate the user who created the attempts by lookin...
[See more](#)

Potential causes

A malicious user is trying to use prompt manipulation t...
[See more](#)

Detected by

Microsoft

Related entities

Account (1)

Azure resource (1)

IP (1) Includes Geo & Threat Intelligence

Supporting evidence events (Preview)

Steps

Now, change to the <https://security.microsoft.com/> portal, navigate to **Investigation & response** -> **Incidents & alerts** -> **Incidents**

Search for **Incident ID 31421**

<https://security.microsoft.com/incident2/31421/overview?tid=0527ecb7-06fb-4769-b324-fd4a3bb865eb>

Note that we have the same Jailbreak alert now contextualized and with the Attack story.

You can navigate the Incident to understand further.

View

The screenshot displays the Microsoft Defender portal interface. The top navigation bar includes the 'Zava' logo, 'Microsoft Defender', a search bar, and user information 'Suresh Dang'. The left sidebar contains a menu with options like Home, Exposure management, Investigation & response, Incidents & alerts, Hunting, Actions & submissions, Partner catalog, Threat intelligence, Assets, Microsoft Sentinel, Identities, Endpoints, Email & collaboration, Cloud apps, and Operational technology. The main content area shows an incident titled 'Multi-stage incident involving Initial access & Privilege escalation involving multiple user...'. It includes a severity indicator (High), status (Active), and a 'Critical asset' tag. A message states: 'This information is limited because of your current permissions. Contact a global administrator to change your permissions.' Below this, there are tabs for 'Attack story', 'Alerts (8)', 'Assets (3)', 'Investigations (0)', 'Evidence and Response (15)', 'Summary', and 'Similar incidents (1)'. The 'Attack story' tab is active, showing a timeline of events: 'Your Azure AI resources have been accessed by an anonymized IP (Preview)', 'Anonymous IP address', and 'A Jailbreak attempt on your Azure AI model deployment was blocked by Prompt Shields'. The 'Incident graph' tab is also visible, showing a network diagram with nodes for 'OAuth Application', 'mdc-d4ai-aoai-1', '195.47.238.87', '7 IPs', and '3 Cloud Applications'. The right sidebar contains 'RELATED THREATS' (Threat Overview Profile: Cloud identity abuse), 'Incident details' (Assigned to: Defender Experts, Incident ID: 31421), 'Classification' (Not set), 'Categories' (Initial access, Execution, Privilege escalation, Suspicious activity), 'First activity' (20 Oct 2025 13:35:34), 'Creation Time' (20 Oct 2025 13:39:17), 'Last activity' (20 Oct 2025 17:07:36), and 'Workspaces'.

Steps

Back to the <https://security.microsoft.com/> portal main console, navigate to Exposure management -> Attack surface -> **Attack paths blade.**

Note there are different attack paths listed.

For illustration purposes open *Device with high severity vulnerabilities allows lateral movement to Azure Open AI* and click on any from the list.

View

The screenshot displays the Microsoft Defender console interface. On the left, the navigation pane shows the 'Attack paths' section highlighted. The main area is titled 'Attack paths' and shows a list of potential attack paths. One path is highlighted: 'Device with high severity vulnerabilities allows lateral movement to Azure Open AI'. The right pane provides a detailed view of this path, including a diagram showing the attack flow from a device to Azure Open AI, and a list of recommendations.

Attack Paths Table:

Attack path name	Entry point	Entry point type	Target	Target type	Target criticality	Type
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	Device	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	cog-204d2d874q04	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	cog-204d2d874q04	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	cog-3nc76c5fmbas	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	cog-3nc76c5fmbas	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	cog-3nc76c5fmbas	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	mdc-44ai-40ai-1	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	cog-204d2d874q04	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	mdc-44ai-40ai-1	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	mdc-44ai-40ai-1	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	cog-204d2d874q04	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	cog-3nc76c5fmbas	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	cog-3nc76c5fmbas	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	cog-204d2d874q04	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	mdc-44ai-40ai-1	Azure OpenAI	None	Hybrid	
Device with high severity vulnerabilities allows lateral movement to Azure Open AI	Device	cog-204d2d874q04	Azure OpenAI	None	Hybrid	

Device with high severity vulnerabilities allows lateral movement to Azure Open AI

Overview Recommendations

alpinexi-u4090 595116751b051a... Smbath Accountability cog-3nc76c5fmbas

Layers
Discovery Source

Description
A device has high severity vulnerabilities The user has high permissions to an Azure OpenAI service Access to Azure OpenAI service resource

Attack Story
1. Attacker can exploit the vulnerabilities and gain control on the device
2. Attacker can use the identity to access, change and list keys of the OpenAI resource
3. Attacker can take over this account, create and manage deployments and keys and access the data

MITRE ATT&CK tactics
Execution
Lateral Movement
Impact


With the Attack path open, review all details provided like:

1. Description
2. Attack story
3. MITRE Att&ck tactics

You can click on **View in Map** to continue exploring the assets

Device with high severity vulnerabilities allows lateral movement to Azure Open AI

Low Azure



Description

A device has high severity vulnerabilities The user has high permissions to an Azure OpenAI service Access to Azure OpenAI service resource

Attack Story

1. Attacker can exploit the vulnerabilities and gain control on the device
2. Attacker can use the identity to access, change and list keys of the OpenAI resource
3. Attacker can take over this account, create and manage deployments and keys and access the data

MITRE ATT&CK® tactics




- > Execution
- > Lateral Movement
- > Impact

Risk factors

Vulnerabilities Lateral movement

View in map

Steps	View
Now, navigate Assets -> AI agents -> AWS Bedrock. Select the agent Note this is the place for Agent inventory and the visibility across different environments. If you click on Open agent page you can continue exploring details of the agent. You can review recommendations and guided steps.	 The screenshot shows the Microsoft Defender AI Agents interface. The left sidebar contains a navigation menu with the following items: Home, Exposure management, Investigation & response, Threat intelligence, Assets, Devices, Identities, Applications, Cloud, AI agents, Microsoft Sentinel, Identities, Endpoints, Email & collaboration, Cloud apps, Cloud security, Operational technology, Security Copilot, Cases, SOC optimization, Reports, Learning hub, and Trials. The 'AI agents' item is highlighted with a red box. The main panel displays the 'AI Agents' section, which includes a list of agents: All (110), Microsoft Foundry (1), Copilot Studio (108), AWS Bedrock (1), and GCP Vertex AI. The 'AWS Bedrock (1)' agent is highlighted with a red box. Below the list, there are two charts: 'Discovered agents' and 'Agents requiring attention', both showing a count of 1. A red arrow points from the 'Open agent page' link in the top right of the main panel to the 'agent-quick-start-88132' agent details panel on the right. The details panel shows the following information: Agent ID: armawsbedrockus-east-1:891377262678:agent/syshdzqzw, Name: agent-quick-start-88132, Creation time: 23 Sept 2025, Attack paths: 0, Risk factors: No risk factors, and Recommendations: 1 Recommendation. A red box highlights the 'Open agent page' link in the top right of the main panel.

Steps	View
	ava Microsoft Defender Search Home Exposure management Investigation & response Threat intelligence Assets Devices Identities Applications Cloud AI agents Microsoft Sentinel Identities Endpoints Email & collaboration Cloud apps Cloud security Operational technology Security Copilot Cases SOC optimization Reports Learning hub Trials AI Agents > agent-quick-start-88132 agent-quick-start-88132 Green Not Prepared Platforms: AWS Bedrock Overview Security recommendations Attack paths Agent info Agent ID arn:aws:bedrock:us-east-1:881377262678:agent/syshdzqzw Name agent-quick-start-88132 Creation time 23 Sept 2025 Attack paths 0 Risk factors No risk factors Map Search AWG Bedrock agents should s... agent-quick-stb... 881377262678 Layers Discovery Source View in map Risk based recommendations 1 Recommendations Critical 0 High 0 2 more View recommendations
References:	- Overview - AI threat protection - Microsoft Defender for Cloud Microsoft Learn - AI and agents Insights Microsoft Security Blog